

```
In [2]: from sklearn.svm import SVC
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import datasets
```

```
In [4]: df=datasets.load_iris(as_frame=True).frame
```

```
In [6]: df.isnull().sum()
```

```
Out[6]: sepal length (cm)      0
sepal width (cm)      0
petal length (cm)      0
petal width (cm)      0
target                  0
dtype: int64
```

```
In [7]: df.shape
```

```
Out[7]: (150, 5)
```

```
In [5]: df.head()
```

```
Out[5]:   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  target
0           5.1              3.5             1.4              0.2            0
1           4.9              3.0             1.4              0.2            0
2           4.7              3.2             1.3              0.2            0
3           4.6              3.1             1.5              0.2            0
4           5.0              3.6             1.4              0.2            0
```

```
In [8]: X=df.drop("target",axis=1)
y=df["target"]
```

```
In [9]: X.head()
```

```
Out[9]:   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0           5.1              3.5             1.4              0.2
1           4.9              3.0             1.4              0.2
2           4.7              3.2             1.3              0.2
3           4.6              3.1             1.5              0.2
4           5.0              3.6             1.4              0.2
```

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)
#stratify=y  Maintains same class distribution in train & test sets
# If dataset has:
# 70% Class A
# 30% Class B
# Then both train & test will preserve this ratio
```

```
In [11]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.transform(X_test)
```

```
In [15]: # model
svc=SVC()
svc.fit(X_train_scaled,y_train)
```

Out[15]:

▼ SVC ⓘ ⓘ

► Parameters

In [16]: `y_pred=svc.predict(X_test_scaled)`

In [17]: `# Evaluate`
`from sklearn.metrics import accuracy_score, classification_report`
`print("accuracy:",accuracy_score(y_test,y_pred))`
`print("classification report:\n",classification_report(y_test,y_pred))`

```
accuracy: 0.9333333333333333
classification report:
      precision    recall   f1-score   support
          0         1.00     1.00     1.00       15
          1         0.88     0.93     0.90       15
          2         0.93     0.87     0.90       15

      accuracy                           0.93        45
   macro avg       0.93     0.93     0.93        45
weighted avg       0.93     0.93     0.93        45
```

In [21]: `# linear kernel`
`svc=SVC(kernel="linear")`
`svc.fit(X_train_scaled,y_train)`
`y_pred=svc.predict(X_test_scaled)`
`print("accuracy:",accuracy_score(y_test,y_pred))`

accuracy: 0.9111111111111111

In [22]: `# Polynomial kernel`
`svc = SVC(kernel="poly")`
`svc.fit(X_train_scaled, y_train)`
`y_pred = svc.predict(X_test_scaled)`
`print("accuracy: ", accuracy_score(y_test, y_pred))`

accuracy: 0.8666666666666667

In [23]: `# Sigmoid kernel`
`svc = SVC(kernel="sigmoid")`
`svc.fit(X_train_scaled, y_train)`
`y_pred = svc.predict(X_test_scaled)`
`print("accuracy: ", accuracy_score(y_test, y_pred))`

accuracy: 0.9111111111111111

C is the Regularization parameter. It controls the trade-off between: ♦ Maximizing margin ♦ Minimizing classification error Small C → Large margin, more misclassification allowed Large C → Small margin, strict classification, less error allowed

In [28]: `C_vals=[0.5,1,2,3,4,5]`
`for c_val in C_vals:`
 `svc=SVC(C=c_val,kernel="rbf")`
 `svc.fit(X_train_scaled, y_train)`
 `y_pred = svc.predict(X_test_scaled)`
 `print("C = ", c_val, "& accuracy: ", accuracy_score(y_test, y_pred))`

cs

```
C = 0.5 & accuracy: 0.9111111111111111
C = 1 & accuracy: 0.9333333333333333
C = 2 & accuracy: 0.9111111111111111
C = 3 & accuracy: 0.9111111111111111
C = 4 & accuracy: 0.9333333333333333
C = 5 & accuracy: 0.9333333333333333
```

C Value	Behavior	Result
Low C (0.01, 0.1)	High regularization	More bias, less variance → underfitting
Medium C (1)	Balanced	Good generalization
High C (10, 100)	Low regularization	Low bias, high variance → overfitting